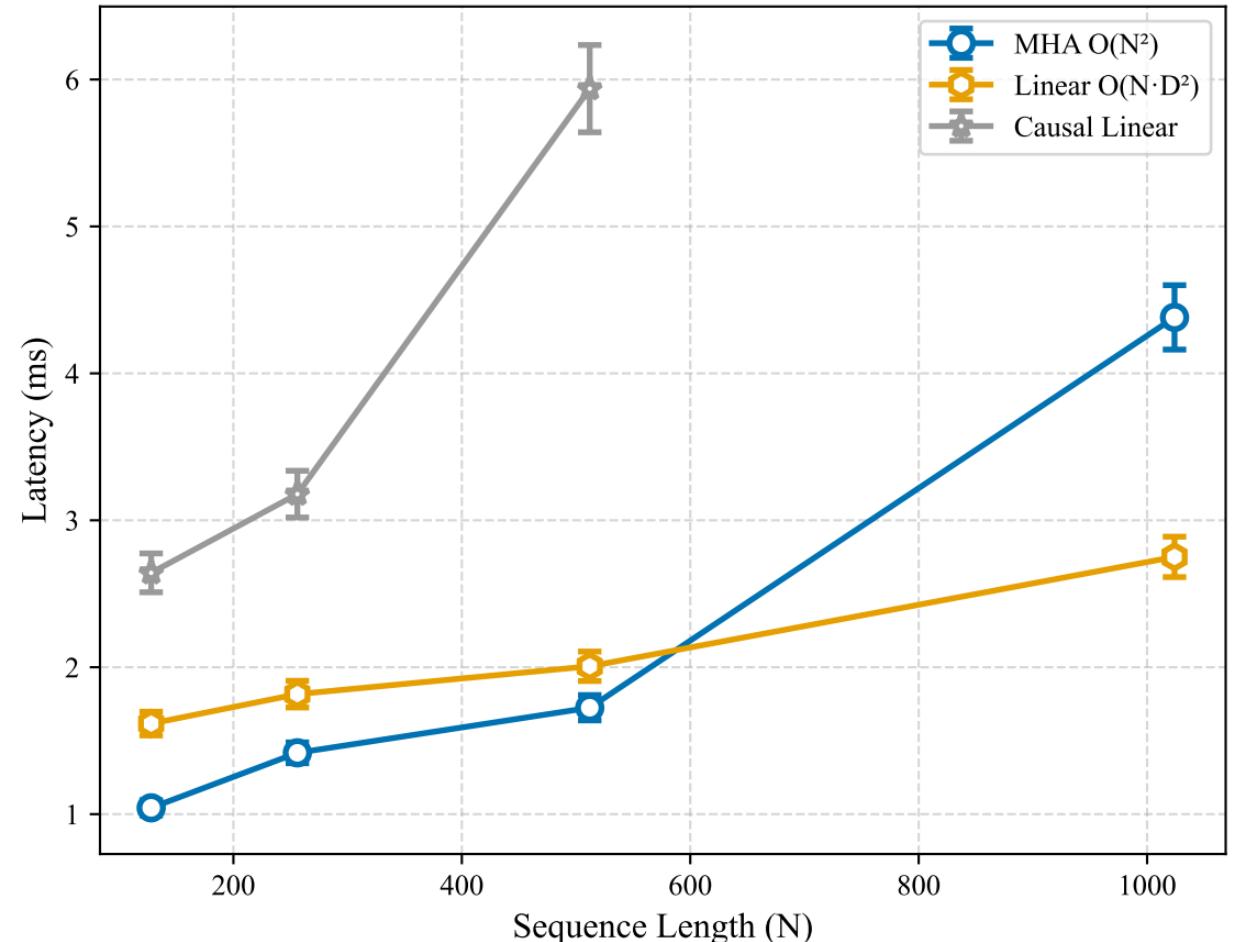


(a) Latency: Quadratic vs Linear Attention



(b) Observed vs Theoretical Scaling

